

Automated Tree Recognition & Green-Space Mapping

Turning a manual, one-by-one task into an **automated** process — from a satellite image to a city tree-map, end to end, for Astana.

AUTHORS Berik Sharipov · Anuar Totin · **Rasul Aidarkhanov** SUPERVISOR S. Satbayev

Y0L0v8x-seg · one capture

AGENDA

What we'll cover

- 01 **Problem & relevance** — why Astana needs this
- 02 **The gap & aim** — off-the-shelf scores ~0 here
- 03 **Objectives, problem statement & methods**
- 04 **Literature review**
- 05 **How we measure success** (the maths)
- 06 **The automated pipeline** (architecture)
- 07 **Data, database & hardware**
- 08 **Four detection engines** — YOLO · Mask R-CNN · DeepForest · SAM 2
- 09 **Why several models** & ensembles
- 10 **Results** on one common test set
- 11 **Canopy** — the web application
- 12 **Discussion & conclusions**

RELEVANCE

Mapping a city's trees — today it's done by hand

Astana is expanding fast. Urban green cover matters for **heat mitigation, irrigation planning and green-space policy** — but to know where the trees are, someone has to **count them one by one** off imagery. That doesn't scale to a whole city.

The real problem isn't "can we beat a benchmark" — it's that **there is no automated process** for this in Astana at all. Existing models are trained on North-American / European data, and we didn't even know if they work here.

So our goal was to **automate the whole process** end-to-end — and to be honest about how well it works.

TODAY — MANUAL

Counting trees by eye, image by image. Slow, not repeatable, never finished for a growing city.

OUR GOAL — AUTOMATED

One repeatable pipeline: point at an area → it finds the trees and puts them on a map. No manual counting.

— WHY WE COULDN'T JUST USE AN EXISTING TOOL

0.012

Box mAP@50

A popular ready-made detector (NEON DeepForest), run as-is on Astana — essentially blind. There was **no off-the-shelf automation** that worked here, so we had to build our own.

AIM

Build an **automated end-to-end system** that recognises trees and maps green space on Astana satellite imagery — and deliver it as a working tool, being honest that the accuracy is modest.

OBJECTIVES

Five tasks

- 1 Build an **annotated Astana dataset** (tiled satellite imagery, instance polygons).
- 2 Train and compare four **detection engines** — YOLOv8-seg, Mask R-CNN, DeepForest, SAM 2.
- 3 **Combine** them through ensembles and test if that helps.
- 4 Evaluate everything on **one common test set** for a fair comparison.
- 5 Wrap it all in an **automated web application** — the actual deliverable.

METHODS

Supervised deep learning · transfer learning from pretrained weights · sliding-window tiling of large images · mAP@50 as the single comparison metric.

— FORMALLY — WHAT GOES IN, WHAT COMES OUT

Input → Output

INPUT

A satellite capture of an area of Astana at zoom 19 (≈ 0.3 m per pixel), cut into overlapping **~640 px RGB tiles** by a sliding window.

Training input also carries human **instance-polygon annotations** — one outline per tree.

OUTPUT

For every detected tree: a **bounding box**, an **instance mask** and a **confidence score**.

Tiles are stitched back, duplicates merged, and the result is shown on a city **map** with per-area counts and **GeoJSON / CSV** export.

```
tiles (RGB) → model f(·) → { box, mask, confidence }per tree → merge & map
```

— LITERATURE REVIEW · 31 PAPERS · 2019–2025

What others report — and on what data

METHOD	AUTHORS · YEAR	DATA	BEST REPORTED METRIC
Mask R-CNN (MCAN)	Lv et al. · 2023	UAV RGB	Detection AP 92.4%
Faster R-CNN / RetinaNet	dos Santos et al. · 2019	UAV RGB	AP 82–93%
YOLOv12m	Abbas & Damaševičius · 2025	RGB satellite	mAP@50 90.8%
DeepLabV3+ / U-Net	Martins / Wang · 2021	Aerial 10–32 cm	F1 91% · IoU 96%
DeepForest (off→fine)	Ventura et al. · 2024	NAIP 60 cm (USA)	F 0.42 → 0.73
SAM / SAM 2	Kirillov / Ravi · 2023–24	1B+ masks	zero-shot segmentation
YOLOv8x-seg (ours)	This work · 2026	Astana satellite · M14	Box mAP@50 0.315

High accuracy — but always on the authors' own region and sensor. Our row is the first measurement of these models on Astana imagery (first, not best).

— THE MATHS, IN PLAIN TERMS

Three simple ideas — one comparison number

1 • IOU — DID THE BOX LAND ON THE TREE?

$$\text{IoU} = \frac{\text{overlap area } (A \cap B)}{\text{combined area } (A \cup B)}$$

How much a predicted box overlaps the real one. **1.0** = perfect, **0** = miss. A tree counts as found when **IoU ≥ 0.5**.

2 • PRECISION & RECALL

Precision — of the trees we flagged, how many were real (few false alarms).

Recall — of all real trees, how many we found: **ours ≈ 30%** at the default setting.

3 • MAP@50 — THE SINGLE SCORE

Sweep the confidence threshold, track precision against recall, and take the **area under that curve** (at IoU ≥ 0.5). One honest number to compare every model the same way — it is the **0.315** on the results slide.


HOW THE TWO DETECTORS WORK

YOLOv8 — one pass: splits the image into a grid and predicts boxes + masks in a single look (fast).

Mask R-CNN — two passes: first proposes regions, then refines them (slower, careful).

— THE AUTOMATED PIPELINE — THE CORE OF THE PROJECT

From a satellite image to a tree-map — automatically



TILING

Large captures are cut into ~640 px tiles with overlap (zoom 19, ≈0.3 m/px); detections are stitched back and de-duplicated.

SWAPPABLE ENGINE

The detection model is just one plug-in step — swap YOLO for Mask R-CNN, DeepForest or an ensemble without touching the rest of the pipeline. ~1 km² at zoom 19 in ≈18 s on a laptop GPU.

— INPUTS & INFRASTRUCTURE

What goes in, and where it runs

DATASET

High-zoom satellite captures of Astana (~0.3 m/px), tiled 640 + 128 overlap. Trees overlap, so we annotate **instance polygons**, not just boxes. A held-out common set of **14 images • 702 labelled trees** scores every model.

DATABASE (SQLITE)

SCAN_SESSIONS

→

SNAPSHOTS

→

RUNS

→

DETECTIONS

One-to-many, **ON DELETE CASCADE**.

HARDWARE

- RTX 4060 8 GB — YOLO + all common-set inference
- RTX 4070 — Mask R-CNN
- RTX 4050 — DeepForest

8 GB VRAM ceiling → batch 2 + mixed precision. Deliberately modest, to prove reproducibility without a cluster.





— FOUR INTERCHANGEABLE ENGINES • ONE PIPELINE


The automation doesn't depend on one model — any of four can drive it.

YOLOv8-seg (Rasul) • Mask R-CNN (Berik) • DeepForest + SAM 2 (Anuar) • plus two ensembles. Each plugs into the same automated pipeline and is scored on the same 14 images.

— MY BRANCH · ~20 EXPERIMENTS

MY CONTRIBUTION

YOLOv8x-seg: from 0.131 to 0.315



Box mAP@50 across every experiment on the common set.

+140%

Box mAP@50 over my own first model (0.131 → **0.315**).

WHAT ACTUALLY MATTERED

- 1 Start from pretrained (COCO) weights
- 2 Tile at the resolution you deploy at
- 3 Keep augmentation moderate

WHAT I ALSO FOUND

Default augmentation beat hand-tuned (top-down trees are rotation-invariant); model size is **U-shaped** (m & x best, l worse); run-to-run noise $\approx \pm 0.03$ — so small gaps aren't real.

— ENGINE #2 · BERIK SHARIPOV

Mask R-CNN — the careful, mask-first engine

A classic **two-stage, region-based** detector. Where YOLO looks once, Mask R-CNN looks twice — propose, then refine — which makes it slower but gives it naturally **crisp instance masks**. In our project it is one of the swappable engines behind the same automated pipeline.

HOW IT WORKS

A **Region Proposal Network** scans each tile and suggests candidate boxes → **RoIAlign** crops every region → three parallel heads output a **box**, a **class** and a per-pixel **mask**. Segmentation is built in, not bolted on, so outlines hug the crown.

HOW BERIK TRAINED IT

Transfer-learned from a **COCO-pretrained** Mask R-CNN and fine-tuned on the same annotated Astana tiles (v2 + v3) for a single **"tree"** class — same train/val split as every other engine, so the comparison is fair. (Exact backbone/epochs to confirm with Berik.)

HOW IT DOES ON ASTANA

Box mAP@50 **0.166**, Mask **0.158**. At confidence 0.5: **precision 0.44** · **recall 0.22** (157 of 702 trees). Strongest on medium, well-separated crowns; like the others, it loses the small and shaded ones.

WHY IT STAYS IN THE APP

Overall below YOLO on our data — but it recovers some **larger trees the others miss**, and its masks are the cleanest. That is exactly why the pipeline keeps it as a user-selectable engine rather than throwing it away.

ENGINE #3 • ANUAR TOTIN

DeepForest + SAM 2 — the forest specialist, re-taught

A two-part engine: **DeepForest** (a tree-crown detector, RetinaNet-style) finds the trees, then **SAM 2** (Segment Anything 2) turns each box into a high-quality mask. It is the engine that produced the **0.012** that started the whole project — and the clearest proof of why domain adaptation matters.

HOW IT WORKS

DeepForest outputs a box per crown; each box becomes a **prompt** for SAM 2, which segments the exact pixels. Two specialist models chained — a detector that knows "tree", and a segmenter that knows "object".

THE DOMAIN PROBLEM

DeepForest is pretrained on **North-American airborne forest** imagery. Pointed at a dry steppe city it was nearly blind — **0.012**. Not a bug, a mismatch: the trees, spacing and sensor all differ from what it had ever seen.

HOW ANUAR FIXED IT


Fine-tuned DeepForest on the annotated Astana data and paired it with SAM 2 masks (confidence 0.3, 400 px patches). Result: **0.012 → 0.146**
Box mAP@50 (Mask **0.134**) — roughly a **×12** jump from re-teaching alone.

WHY IT STAYS IN THE APP

SAM 2 gives the **best general-purpose masks** of any engine, and the off-the-shelf→fine-tuned story is the project's core lesson. It still struggles with tiny, dense urban trees — so it's offered alongside the others, not instead of them.

— THE MOST USEFUL LESSON — AUTOMATION, NOT ONE MODEL

Same garden, eight engines, eight answers



The very same Botanical Garden, scanned by **eight detection engines** — each outlines a **different set of trees** and reports a different count (**672-819**). None is simply "right".


That's the whole design rationale: the pipeline lets the user **switch engines** and combine them — **Weighted Box Fusion** + a **cross-YOLO 4-vote** that drops single-model false positives (e.g. stadium roofs).

— HEAD-TO-HEAD • SAME 14 IMAGES • 702 TREES

One common test set, every model

MODEL	BOX MAP@50	MASK MAP@50
YOLOv8x-seg v4 (ours)	0.315	0.289
YOLOv8x-seg v3	0.287	0.263
Mask R-CNN	0.166	0.158
DeepForest + SAM 2	0.146	0.134
YOLO v1 baseline	0.131	0.134
NEON DeepForest (off-the-shelf)	0.012	—

Champion operating point (conf 0.25): **precision 0.52 • recall 0.29** (222 of 755 trees). Margins between top models are small – partly within run-to-run noise.



We read this as "these engines are in the same ballpark on Astana," **not** "one clearly wins." What matters for the project: **any of them can drive the automated pipeline.**


tree 0.3 tree 0.3 tree 0.5

WHAT IT ACTUALLY LOOKS LIKE

Detections vs. real crowns

Predictions line up with real tree crowns in many places. But honestly — it **doesn't find every tree**: at the app's default confidence it finds about **30%** of labelled trees (more if looser, fewer if stricter), and per image it swings from almost none to nearly all.


tree 0.4 tree 0.4 tree 0.6



— THE DELIVERABLE

MY CONTRIBUTION

Canopy — pick an area, pick a model, scan




ONE FLOW

Select an area of Astana → the app tiles it, runs the chosen model, and draws every detected tree on the map — with a confidence filter and Point / Box / Polygon / Heat layers.

FREE CHOICE OF MODEL

YOLOv8 (7 sizes) · DeepForest ± SAM 2 · Mask R-CNN · Ensemble — because different models suit different images.

Counts shown in screenshots come from repeated demo scans — **not** a full city census.




Region scan over a rectangle or polygon, with live progress streaming. The user chooses the model and size before running — the multi-model idea, made concrete in the UI.

YOLO · DEEFOREST · MASK R-CNN · ENSEMBLE

— THE PIPELINE RUNNING END-TO-END

MY CONTRIBUTION

Live demo — automated scan of the Botanical Garden



WHAT YOU'RE SEEING

The v4 engine scanning the Botanical Garden end-to-end in Canopy: **tile** → **detect** → **map**, every tree drawn live.

WHY THIS SCENE

A dense, well-defined planting — the clearest case. Honest reminder: on sparser scenes recall drops.

FULL DEMO VIDEO

Plays here automatically (~1 min) — or open the full recording:

https://drive.google.com/file/d/10ceaLC4yMEgLtIThVR8t7_B4VRRPNaNf/view

HONEST ASSESSMENT

What we'd say up front

LIMITATIONS

- Small training set → modest accuracy (0.315 mAP, ≈30% recall)
- Domain shift: training captures vs. live map tiles
- Misses small & heavily shaded trees

BUT ALSO

- Off-the-shelf fails here (0.012) — our models do meaningfully better
- A fair, reproducible comparison on local data
- A working tool, not just a notebook

None of this is hidden — it's the honest state of a student project, and each limit improves with more annotated data.

 OBJECTIVES → OUTCOMES

Every task, addressed

- ✓ **An automated pipeline** — capture → tile → detect → map → export, with no manual counting.
- ✓ **Working application** — Canopy, deployed, with a free choice of detection engine.
- ✓ **Dataset** — annotated, tiled Astana imagery with a held-out 14-image / 702-tree common set.
- ✓ **Four engines compared** + two ensembles, all scored identically on one set.
- ✓ **Fair evaluation** — first such measurement for Astana.

The headline isn't a score — it's that a **manual task is now an automated process**. Accuracy is modest (and honestly so: ready-made tools score ~0 here, ours do meaningfully better), but the **process** is the contribution.

MY CONTRIBUTION (RASUL)

- 1 The whole **automated pipeline & Canopy app** (FastAPI + React + SQLite) — the actual deliverable
- 2 The **YOLOv8 engine** — ~20 experiments, +140% over my baseline
- 3 The **cross-YOLO vote** ensemble & the common-set evaluation

FUTURE WORK

- 1 More annotated data
- 2 Train on the **same imagery we run on** (close the domain gap)
- 3 Better handling of small trees · city-wide canopy analytics

Scan area Polygon

ASTANA - CANOPY AGGREGATE

51 348 trees

across 14 scans - 72 snapshots

58 3668 13 749
runs avg / scan visible

32% 4.0 m
avg confidence avg crown

FILTERS
Min confidence 0%
High Med Low

THANK YOU
RECENT SCANS Manage

Thank you — we'll gladly
answer your questions.

TEAM Sharipov • Totin • Aidarkhanov SUPERVISOR S. Satbayev